

# 1. Análise Exploratória de Dados

```
# Importando bibliotecas necessárias
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
import os

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

from sklearn.tree import export_graphviz
import graphviz #
https://stackoverflow.com/questions/33433274/anaconda-graphviz-cant-import-after-installation

# Definindo caminho dos dados
dir_path = './CMAPSSData/'
train_file = 'train_FD003.txt'
test_file = 'test_FD003.txt'

# Definindo o nome das colunas para facilitar a exploração dos dados
index_names = ['unidade', 'ciclo_tempo']
setting_names = ['config_1', 'config_2', 'config_3']
sensor_names = ['s_{}'.format(i+1) for i in range(0,21)]
col_names = index_names + setting_names + sensor_names

# Lendo os dados
train = pd.read_csv((dir_path+train_file), sep='\s+', header=None,
                    names=col_names)
test = pd.read_csv((dir_path+test_file), sep='\s+', header=None,
                  names=col_names)
y_test = pd.read_csv((dir_path+'RUL_FD003.txt'), sep='\s+',
                    header=None,
                    names=['RemainingUsefulLife'])

# Analisar as primeiras linhas da nossa base de dados
print(train.shape)
train.head()
```

(24720, 26)

	unidade	ciclo_tempo	config_1	config_2	config_3	s_1	s_2
0	1	1	-0.0005	0.0004	100.0	518.67	642.36
1	1	2	0.0008	-0.0003	100.0	518.67	642.50

2	1	3	-0.0014	-0.0002	100.0	518.67	642.18
3	1	4	-0.0020	0.0001	100.0	518.67	642.92
4	1	5	0.0016	0.0000	100.0	518.67	641.68

	s_3	s_4	s_5	...	s_12	s_13	s_14	s_15
s_16	s_17	\						
0	1583.23	1396.84	14.62	...	522.31	2388.01	8145.32	8.4246
0.03	391							
1	1584.69	1396.89	14.62	...	522.42	2388.03	8152.85	8.4403
0.03	392							
2	1582.35	1405.61	14.62	...	522.03	2388.00	8150.17	8.3901
0.03	391							
3	1585.61	1392.27	14.62	...	522.49	2388.08	8146.56	8.3878
0.03	392							
4	1588.63	1397.65	14.62	...	522.58	2388.03	8147.80	8.3869
0.03	392							

	s_18	s_19	s_20	s_21
0	2388	100.0	39.11	23.3537
1	2388	100.0	38.99	23.4491
2	2388	100.0	38.85	23.3669
3	2388	100.0	38.96	23.2951
4	2388	100.0	39.14	23.4583

[5 rows x 26 columns]

*# Analisando os ciclos de tempo*

train[index\_names].groupby('unidade').max().describe()

	ciclo_tempo
count	100.00000
mean	247.20000
std	86.48384
min	145.00000
25%	189.75000
50%	220.50000
75%	279.75000
max	525.00000

*# Inspeccionando alguns dados estatisticos do conjunto*

*# Sensores com desvio padrao proximos de 0 nao agregam informacao ao modelo*

train[sensor\_names].describe().transpose()

	count	mean	std	min	25%
50%	\				
s_1	24720.0	518.670000	6.684921e-11	518.6700	518.6700

518.6700					
s_2	24720.0	642.457858	5.230311e-01	640.8400	642.0800
642.4000					
s_3	24720.0	1588.079175	6.810418e+00	1564.3000	1583.2800
1587.5200					
s_4	24720.0	1404.471212	9.773178e+00	1377.0600	1397.1875
1402.9100					
s_5	24720.0	14.620000	3.602525e-12	14.6200	14.6200
14.6200					
s_6	24720.0	21.595841	1.811600e-02	21.4500	21.5800
21.6000					
s_7	24720.0	555.143808	3.437343e+00	549.6100	553.1100
554.0500					
s_8	24720.0	2388.071555	1.582849e-01	2386.9000	2388.0000
2388.0700					
s_9	24720.0	9064.110809	1.998029e+01	9017.9800	9051.9200
9060.0100					
s_10	24720.0	1.301232	3.484849e-03	1.2900	1.3000
1.3000					
s_11	24720.0	47.415707	3.000742e-01	46.6900	47.1900
47.3600					
s_12	24720.0	523.050873	3.255314e+00	517.7700	521.1500
521.9800					
s_13	24720.0	2388.071643	1.581207e-01	2386.9300	2388.0100
2388.0700					
s_14	24720.0	8144.202916	1.650412e+01	8099.6800	8134.5100
8141.2000					
s_15	24720.0	8.396176	6.051161e-02	8.1563	8.3606
8.3983					
s_16	24720.0	0.030000	1.750371e-14	0.0300	0.0300
0.0300					
s_17	24720.0	392.566545	1.761459e+00	388.0000	391.0000
392.0000					
s_18	24720.0	2388.000000	0.000000e+00	2388.0000	2388.0000
2388.0000					
s_19	24720.0	100.000000	0.000000e+00	100.0000	100.0000
100.0000					
s_20	24720.0	38.988552	2.488647e-01	38.1700	38.8300
38.9900					
s_21	24720.0	23.393024	1.492338e-01	22.8726	23.2962
23.3916					

	75%	max
s_1	518.6700	518.6700
s_2	642.7900	645.1100
s_3	1592.4125	1615.3900
s_4	1410.6000	1441.1600
s_5	14.6200	14.6200
s_6	21.6100	21.6100

s_7	556.0400	570.4900
s_8	2388.1400	2388.6000
s_9	9070.0925	9234.3500
s_10	1.3000	1.3200
s_11	47.6000	48.4400
s_12	523.8400	537.4000
s_13	2388.1400	2388.6100
s_14	8149.2300	8290.5500
s_15	8.4370	8.5705
s_16	0.0300	0.0300
s_17	394.0000	399.0000
s_18	2388.0000	2388.0000
s_19	100.0000	100.0000
s_20	39.1400	39.8500
s_21	23.4833	23.9505

```
def add_RUL(df):
    # Obter o numero total de ciclos para cada unidade
    grouped_by_unit = df.groupby(by="unidade")
    max_cycle = grouped_by_unit["ciclo_tempo"].max()

    # Mesclar o valor do ciclo maximo no dataframe de origem
    result_frame = df.merge(max_cycle.to_frame(name='ciclo_max'),
                             left_on='unidade', right_index=True)


    # Calcular o RUL para cada linha
    remaining_useful_life = result_frame["ciclo_max"] -
result_frame["ciclo_tempo"]
    result_frame["RUL"] = remaining_useful_life

    # Remover o valor do ciclo maximo, que nao e mais necessario
    result_frame = result_frame.drop("ciclo_max", axis=1)
    return result_frame
```

```
train = add_RUL(train)
train[index_names+['RUL']].head()
```




	unidade	ciclo_tempo	RUL
0	1	1	258
1	1	2	257
2	1	3	256
3	1	4	255
4	1	5	254




```
df_max_rul = train[['unidade',
'RUL']].groupby('unidade').max().reset_index()
df_max_rul['RUL'].hist(bins=15, figsize=(15,7))
plt.xlabel('RUL')
plt.ylabel('Frequência')
plt.show()
```






## 2. Plotando sinais




```
def plot_sinal(df, signal_name):  
    plt.figure(figsize=(13,5))  
    for i in df['unidade'].unique():  
        if (i % 10 == 0):  
            plt.plot('RUL', signal_name,  
                     data=df[df['unidade']==i])  
    plt.xlim(250, 0) # Inverte o eixo x para ir de 250 ate 0  
    plt.xticks(np.arange(0, 300, 25))  
    plt.ylabel(signal_name)  
    plt.xlabel('RUL')  
    plt.show()  
  
for sensor in sensor_names:  
    plot_sinal(train, sensor)
```


























### 3. Modelo de Referência

```
# Preparando os dados

# Removendo os sensores desnecessarios
drop_sensors = ['s_1','s_5','s_16','s_18','s_19'] # Vou manter o s_6
e s_10 por enquanto
drop_labels = index_names+setting_names+drop_sensors
remaining_sensors = ['s_2', 's_3', 's_4', 's_6', 's_7', 's_8', 's_9',
's_10',
                    's_11', 's_12', 's_13', 's_14', 's_15', 's_17', 's_20', 's_21']

# Separando o RUL do grupo de treino
X_train = train.drop(drop_labels, axis=1)
y_train = X_train.pop('RUL')
y_train_clipped = y_train.clip(upper=125) # Cortando o RUL

# O valor real do RUL esta no ultimo ciclo de tempo de cada unidade
# Reorganizando o grupo de teste com esse ajuste
X_test =
test.groupby('unidade').last().reset_index().drop(drop_labels, axis=1)

# Criando a funcao de avaliacao
def avaliar(y_verdadeiro, y_calculado, label='teste'):
    mse = mean_squared_error(y_verdadeiro, y_calculado)
    rmse = np.sqrt(mse)
    variancia = r2_score(y_verdadeiro, y_calculado)
    print('conjunto de {} -> RMSE:{}, R2:{}'.format(label, rmse,
variancia))

# Criando e ajustando o modelo
lm = LinearRegression()
lm.fit(X_train, y_train)

# Testando e avaliando o modelo treinado
y_hat_train = lm.predict(X_train)
avaliar(y_train, y_hat_train, 'treino')

y_hat_test = lm.predict(X_test)
avaliar(y_test, y_hat_test)

conjunto de treino -> RMSE:63.494471921833146, R2:0.5873657567234036
conjunto de teste -> RMSE:57.05730530062128, R2:-0.8998253079767999
```

### 4. Floresta Aleatória - Random Forest

```
# Criando e ajustando o regressor de floresta aleatória sem qualquer
alteracao
```

```

rf = RandomForestRegressor(n_estimators=100, max_features="sqrt",
random_state=42)
rf.fit(X_train, y_train_clipped)

# Testando e avaliando o modelo treinado
y_hat_train = rf.predict(X_train)
avaliar(y_train_clipped, y_hat_train, 'treino')

y_hat_test = rf.predict(X_test)
avaliar(y_test, y_hat_test)

conjunto de treino -> RMSE:5.9199939580022525, R2:0.9787661901585051
conjunto de teste -> RMSE:21.05308450085165, R2:0.7413439613827657

# Analisando a forma de uma das arvores da floresta
print(rf.estimators_[5].tree_.max_depth) # Quantos nos no caminho
mais longo
rf.estimators_[5].tree_.n_node_samples # Quantos samples nos
ultimos nos

33

array([15616, 11694, 7793, ..., 1, 1, 4], dtype=int64)

# Refazendo o modelo com alguns ajustes
rf = RandomForestRegressor(n_estimators=100, max_features="sqrt",
random_state=42,
max_depth=8, min_samples_leaf=50)
rf.fit(X_train, y_train_clipped)

# Testando e avaliando o modelo treinado
y_hat_train = rf.predict(X_train)
avaliar(y_train_clipped, y_hat_train, 'treino')

y_hat_test = rf.predict(X_test)
avaliar(y_test, y_hat_test)

conjunto de treino -> RMSE:15.706704198492831, R2:0.8505294865338602
conjunto de teste -> RMSE:20.994958823842456, R2:0.7427702419664686

```

## 5. Enxergando a RF

```

def export_rf_visual(estimator):
    # Criando representacao da arvore de decisao
    dot_data = export_graphviz(estimator,
out_file=None,
feature_names=X_train.columns,
filled=True,
special_characters=True,

```

```
        rotate=True)  
  
    graph = graphviz.Source(dot_data)  
    return graph  
  
# Exibindo no notebook  
export_rf_visual(rf.estimators_[5])
```





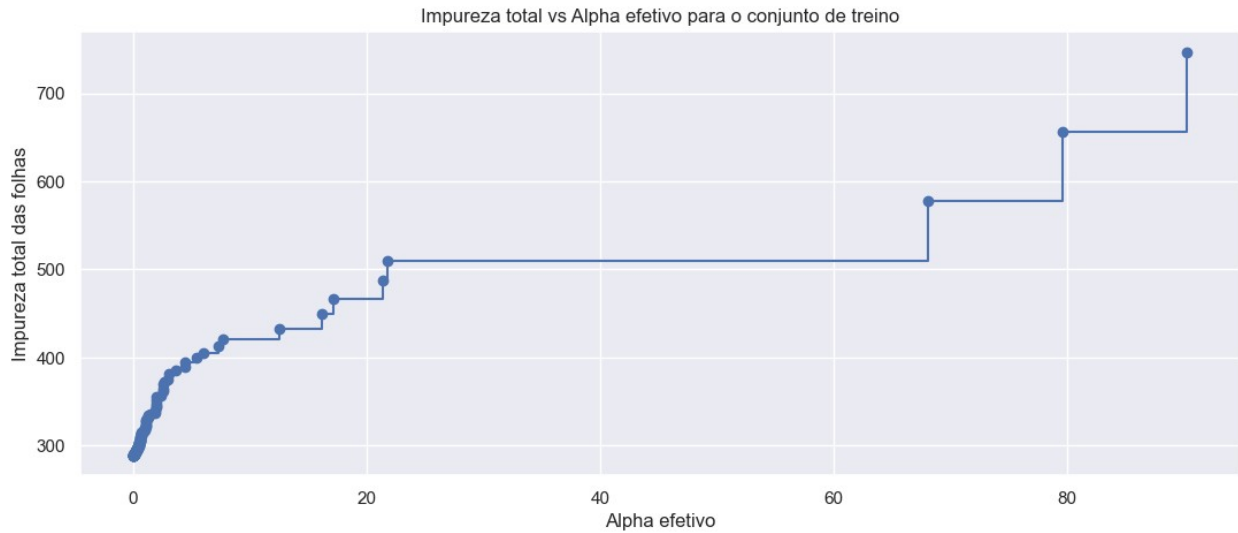
## 6. Hiperparâmetros

```
rf.get_params()

{'bootstrap': True,
 'ccp_alpha': 0.0,
 'criterion': 'squared_error',
 'max_depth': 8,
 'max_features': 'sqrt',
 'max_leaf_nodes': None,
 'max_samples': None,
 'min_impurity_decrease': 0.0,
 'min_samples_leaf': 50,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'n_estimators': 100,
 'n_jobs': None,
 'oob_score': False,
 'random_state': 42,
 'verbose': 0,
 'warm_start': False}

# Obtendo ccp_alphas
path = rf.estimators_[5].cost_complexity_pruning_path(X_train,
y_train_clipped)
ccp_alphas, impurities = path.ccp_alphas, path.impurities

fig, ax = plt.subplots(figsize=(13,5))
ax.plot(ccp_alphas[:-1], impurities[:-1], marker='o',
drawstyle="steps-post")
ax.set_xlabel("Alpha efetivo")
ax.set_ylabel("Impureza total das folhas")
ax.set_title("Impureza total vs Alpha efetivo para o conjunto de
treino")
plt.show()
```



```
# zoom in
fig, ax = plt.subplots(figsize=(13,5))
ax.plot(ccp_alphas[:-10], impurities[:-10], marker='o',
drawstyle="steps-post")
ax.set_xlabel("Alpha efetivo")
ax.set_ylabel("Impureza total das folhas")
ax.set_title("Impureza total vs Alpha efetivo para o conjunto de
treino")
plt.show()
```



```
rf_dict = {
'id_node': list(range(rf.estimators_[5].tree_.node_count)),
'impurity': rf.estimators_[5].tree_.impurity,
'samples': rf.estimators_[5].tree_.n_node_samples,
'id_left_child': rf.estimators_[5].tree_.children_left,
```

```

    'id_right_child': rf.estimators_[5].tree_.children_right
}

```

```

impurity_df = pd.DataFrame(rf_dict)
print(impurity_df.shape)
impurity_df.head(10)

```

```
(227, 5)
```

	id_node	impurity	samples	id_left_child	id_right_child
0	0	1644.689840	15616	1	128
1	1	609.414829	11694	2	71
2	2	247.609863	7793	3	52
3	3	156.037153	6372	4	29
4	4	97.827624	4473	5	20
5	5	19.630575	1627	6	13
6	6	3.844416	765	7	10
7	7	10.192376	257	8	9
8	8	1.085113	143	-1	-1
9	9	21.980353	114	-1	-1

```

# Calculando o min_impurity_decrease

```

```

impurity_df['impurity_decrease'] = np.nan
samples_total = rf.estimators_[5].tree_.node_count

```

```

for idx in impurity_df.index[1:]: # Pulando o primeiro no, ja que n
ha divisoes antes dele

```

```

    if impurity_df.iloc[idx]['id_left_child'] == -1:
        continue # n eh possivel calcular o decaimento da impureza
para folhas, ja que n ha divisoes apos elas

```

```

    else:
        impurity_P, samples_P = impurity_df.iloc[idx][['impurity',
'samples']]
        id_L, id_R = impurity_df.iloc[idx][['id_left_child',
'id_right_child']].astype(int)
        impurity_L, samples_L = impurity_df.iloc[id_L][['impurity',
'samples']]
        impurity_R, samples_R = impurity_df.iloc[id_R][['impurity',
'samples']]

```

```

        impurity_decrease = samples_P / samples_total * (
            impurity_P - samples_R / samples_P * impurity_R -
            samples_L / samples_P * impurity_L
        )

```

```

        impurity_df.at[idx, 'impurity_decrease'] = impurity_decrease

```

```

# Histograma do decaimento de impureza

```

```


impurity_df['impurity_decrease'].plot(kind='hist', bins=50)

```

```

<Axes: ylabel='Frequency'>

```




```
# Dados estatísticos
impurity_df['impurity_decrease'].describe()

count      112.000000
mean       282.515265
std        994.407842
min        -0.506232
25%         14.589094
50%         47.024819
75%        129.095879
max        7639.300570
Name: impurity_decrease, dtype: float64

# Zoom no histograma
impurity_df.loc[impurity_df['impurity_decrease'] < 14.59,
'impurity_decrease'].plot(kind='hist', bins=15)

<Axes: ylabel='Frequency'>
```



## 7. Procura aleatória - Randomsearch

```
# Definindo os limites de cada parametro para a randomsearch
max_depth = [None] + list(range(3, 34, 3))
min_samples_leaf = list(range(1, 102, 10))
min_impurity_decrease = list(np.arange(0, 147)/10)
ccp_alpha = list(np.round(np.linspace(0, 2, 81), decimals=3))

parameters = {
    'max_depth': max_depth,
    'min_samples_leaf': min_samples_leaf,
    'min_impurity_decrease': min_impurity_decrease,
    'ccp_alpha': ccp_alpha
}

tuning_options = len(max_depth) * len(min_samples_leaf) *
len(min_impurity_decrease) * len(ccp_alpha)
print(tuning_options)

1571724

# Criando um conjunto de validação e evitando o vazamento dos dados
com o uso do GroupKFold
```

```

from sklearn.model_selection import RandomizedSearchCV, GroupKFold

ITERATIONS = 300

rf = RandomForestRegressor(n_estimators=100, max_features="sqrt",
random_state=42)
gkf = GroupKFold(n_splits=3)

regressor = RandomizedSearchCV(rf,
                                parameters,
                                cv = gkf.split(train,
groups=train['unidade']),
                                verbose=2,
                                error_score='raise',
                                n_iter=ITERATIONS,
                                n_jobs=-2,
                                # scoring
https://scikit-learn.org/stable/modules/model\_evaluation.html#scoring-
parameter
                                scoring= 'neg_root_mean_squared_error')

regressor.fit(X_train, y_train_clipped)

Fitting 3 folds for each of 300 candidates, totalling 900 fits

RandomizedSearchCV(cv=<generator object _BaseKFold.split at
0x000000189DD648740>,
                    error_score='raise',
                    estimator=RandomForestRegressor(max_features='sqrt',
                                                    random_state=42),
                    n_iter=300, n_jobs=-2,
                    param_distributions={'ccp_alpha': [0.0, 0.025,
0.05, 0.075,
                                                    0.1, 0.125,
0.15, 0.175,
                                                    0.2, 0.225,
0.25, 0.275,
                                                    0.3, 0.325,
0.35, 0.375,
                                                    0.4, 0.425,
0.45, 0.475,
                                                    0.5, 0.525,
0.55, 0.575,
                                                    0.6, 0.625,
0.65, 0.675,
                                                    0.7,
0.725, ...],
                    'max_depth': [None, 3, 6, 9,

```

```

12, 15, 18,
21, 24, 27, 30,
33],
'min_impurity_decrease': [0.0,
0.1, 0.2,
0.3,
0.4, 0.5,
0.6,
0.7, 0.8,
0.9,
1.0, 1.1,
1.2,
1.3, 1.4,
1.5,
1.6, 1.7,
1.8,
1.9, 2.0,
2.1,
2.2, 2.3,
2.4,
2.5, 2.6,
2.7,
2.8, 2.9, ...],
'min_samples_leaf': [1, 11,
21, 31, 41,
51, 61,
71, 81, 91,
101]},
scoring='neg_root_mean_squared_error', verbose=2)

resultado = pd.DataFrame(regressor.cv_results_)
columns = ['param_min_samples_leaf', 'param_min_impurity_decrease',
'param_max_depth', 'param_ccp_alpha',
'mean_test_score', 'std_test_score', 'rank_test_score']
resultado[columns].sort_values('mean_test_score', ascending=False)

param_min_samples_leaf param_min_impurity_decrease param_max_depth
\
26 21 0.1 21
277 31 0.2 33
299 1 0.0 15
123 31 0.8 9
4 11 1.2 24
.. ... ...

```

295	101	8.6	3
84	31	9.1	3
223	71	10.4	3
3	11	11.4	3
212	71	13.7	3

	param_ccp_alpha	mean_test_score	std_test_score	rank_test_score
26	0.475	-16.669104	0.316564	1
277	0.35	-16.744735	0.332176	2
299	0.875	-16.813201	0.281859	3
123	0.175	-16.942223	0.298950	4
4	0.85	-17.041977	0.291904	5
..	...	...	...	...
295	0.925	-19.089300	0.524540	296
84	1.275	-19.090969	0.523144	297
223	0.85	-19.098670	0.521663	298
3	0.475	-19.106038	0.521172	299
212	1.9	-19.122912	0.513109	300

[300 rows x 7 columns]

*# Melhores parametros*

`print(regressor.best_params_)`

`print(regressor.best_score_)`

```
{'min_samples_leaf': 21, 'min_impurity_decrease': 0.1, 'max_depth':
21, 'ccp_alpha': 0.475}
-16.669103655003024
```

## 8. Modelo Final

*# Usando os melhores parametros para treinar o modelo final*

`rf = RandomForestRegressor(n_estimators=100, max_features="sqrt",`



```
random_state=42,
min_impurity_decrease=0.0, min_samples_leaf=11,
max_depth=15, ccp_alpha=0.125)
rf.fit(X_train, y_train_clipped)

# Testando e avaliando o modelo
y_hat_train = rf.predict(X_train)
avaliar(y_train_clipped, y_hat_train, 'treino')

y_hat_test = rf.predict(X_test)
avaliar(y_test, y_hat_test)

conjunto de treino -> RMSE:13.95446880579081, R2:0.8820190156933622
conjunto de teste -> RMSE:20.61288923394374, R2:0.7520472702746352
```